

An LLVM based NPU compiler on RISC-V

In Optical Computing

Jonathan-cs chen June 02, 2023

My jobs in LT

C++ toolchain

- Build gnu and llvm toolchain for RISCV from open source
 - -> LT is able to negotiate with Andes for better price
- Survey open sources TVM, SYCL/OpenCL
 - Reference
- Survey and evaluate vendors' SW and HW in AI and data science applications for optical computing
- Lead the SW development for Auroa HW product and program the compiler backend myself

Optical Computing

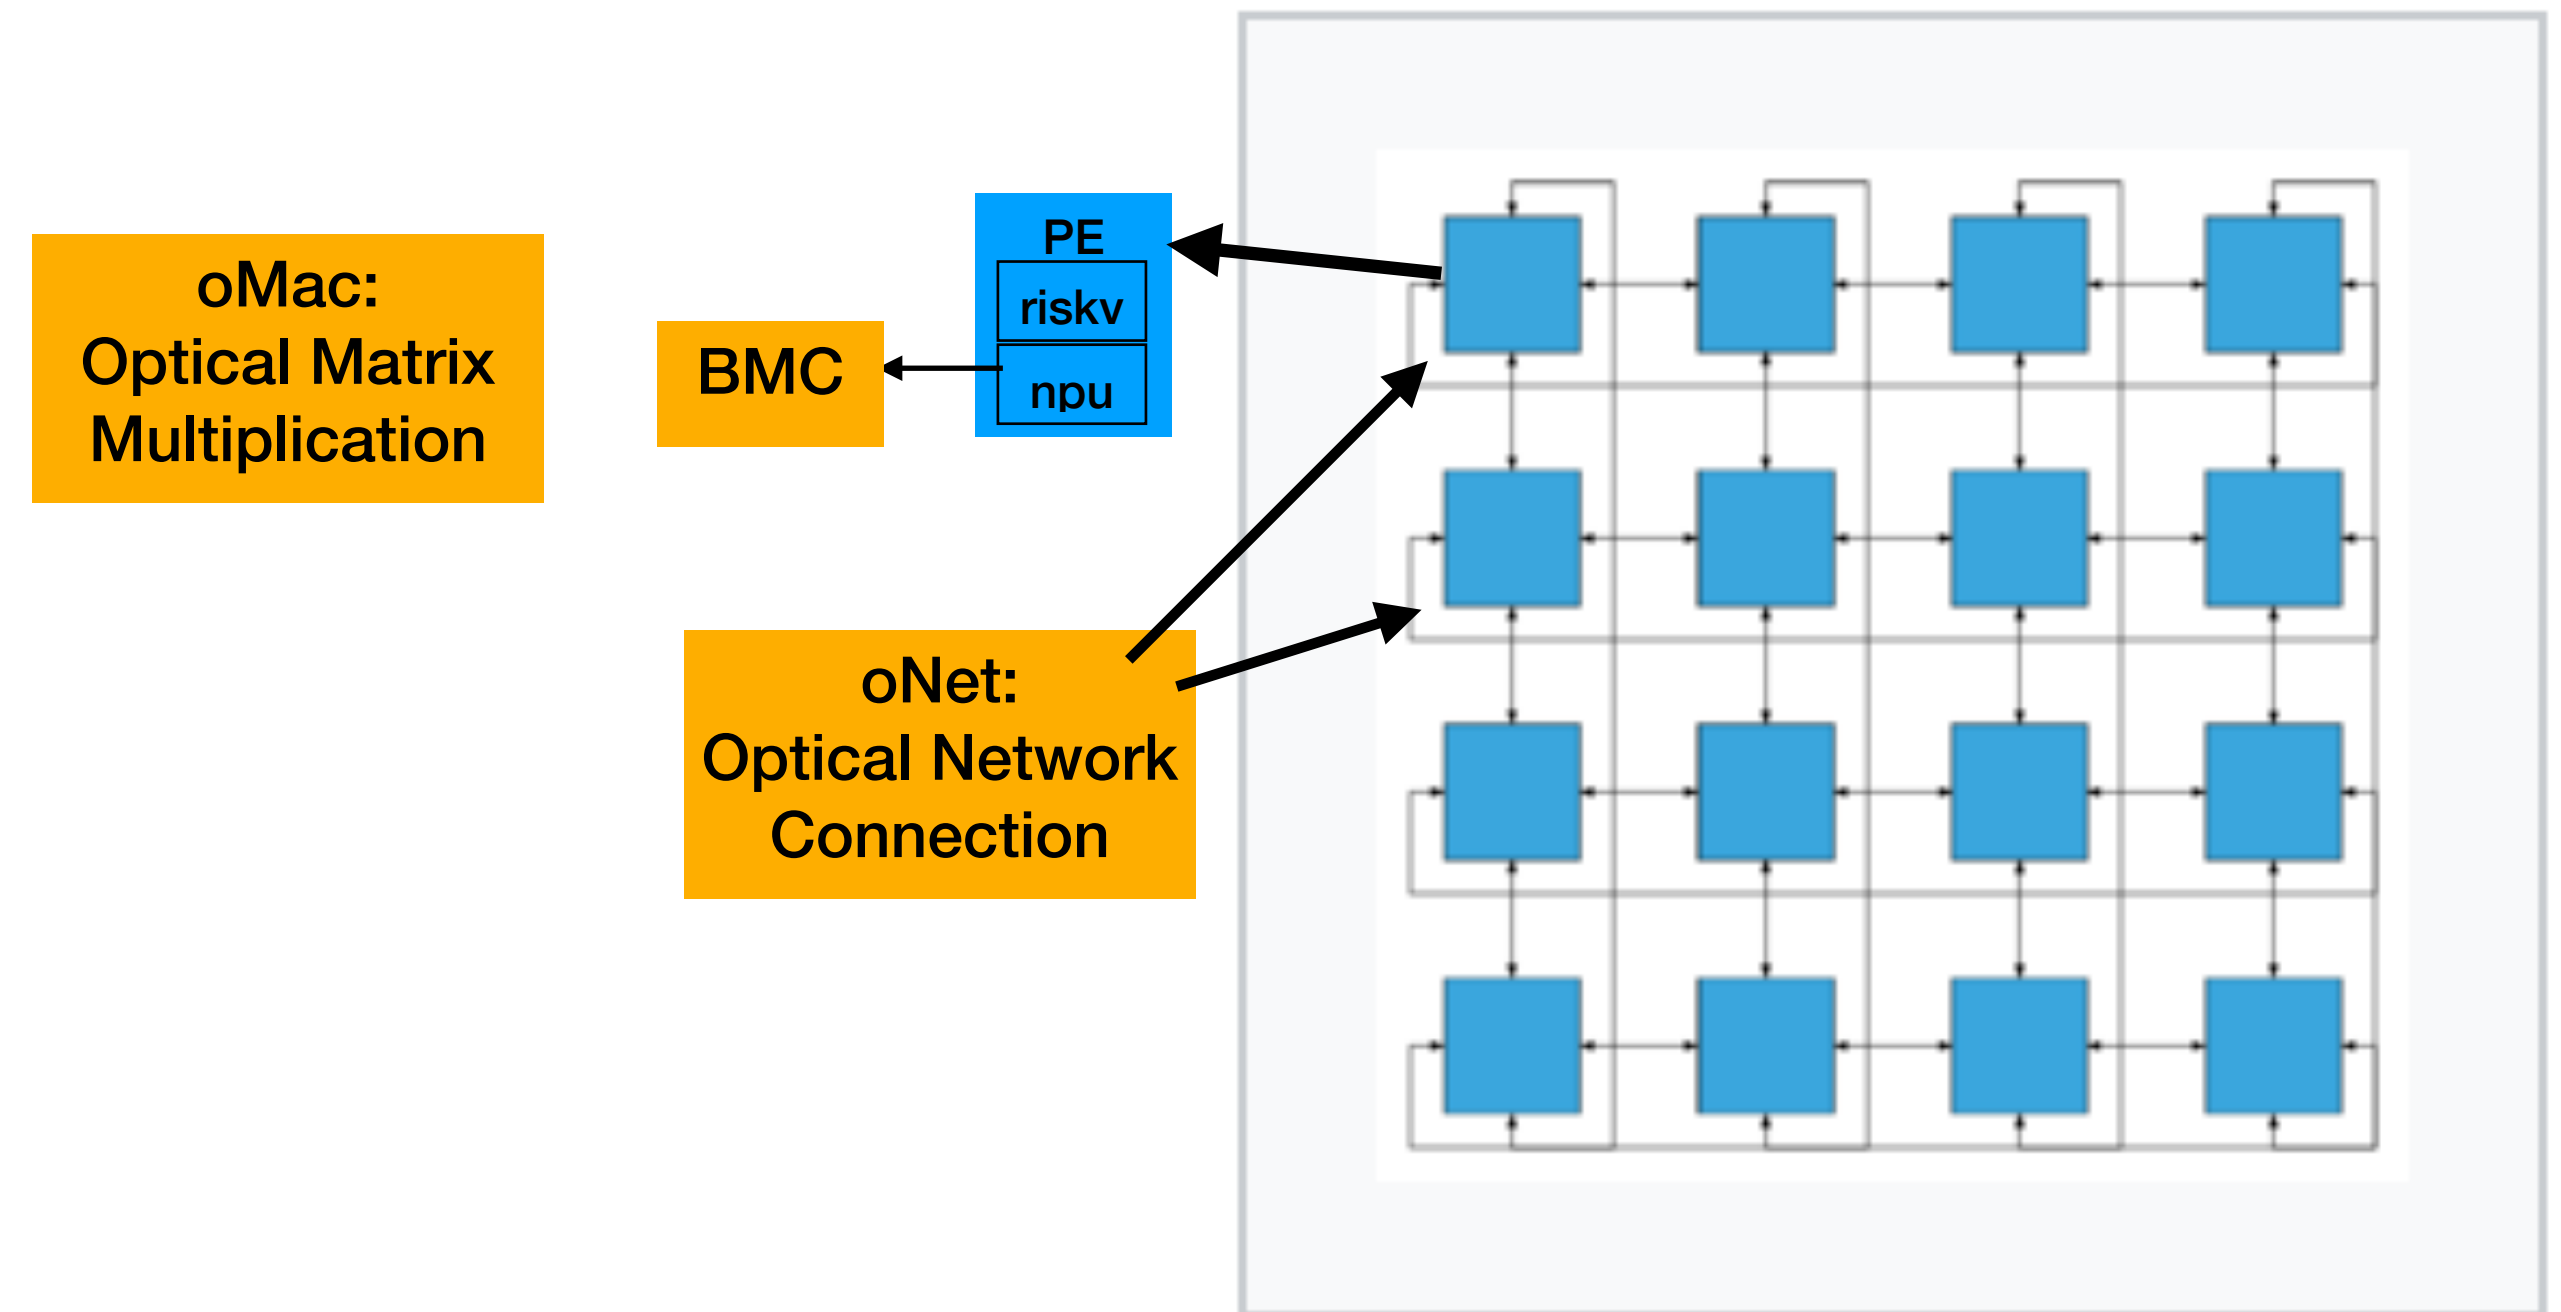
Features and Application

- Multiplexing signals:
 - Using different wavelengths of light for each signal means multiple signals can be sent down the same waveguide at the same time. Electric logic gate has 1 pair of bits only, for instance: adder, and, or gates.
 - Theory: 1,000 channels. Practice: 8 , 10, ... (increased as laser-sensor-tech advanced).
- Application: (ref)
 - Computing: Use 10 pair of bits at same time then get exactly computing.
 - Simulation: 1000個電訊號透過波導集成後餵給矩陣計算器，1000組顏色波長間較近，波導間有干涉、不準，但可做1000組矩陣計算的模擬。
 - Manufacture: From 2010, every big semi-conductor manufactory has their production-line for Optical Computing.

Optical Computing

Architecture : Reference

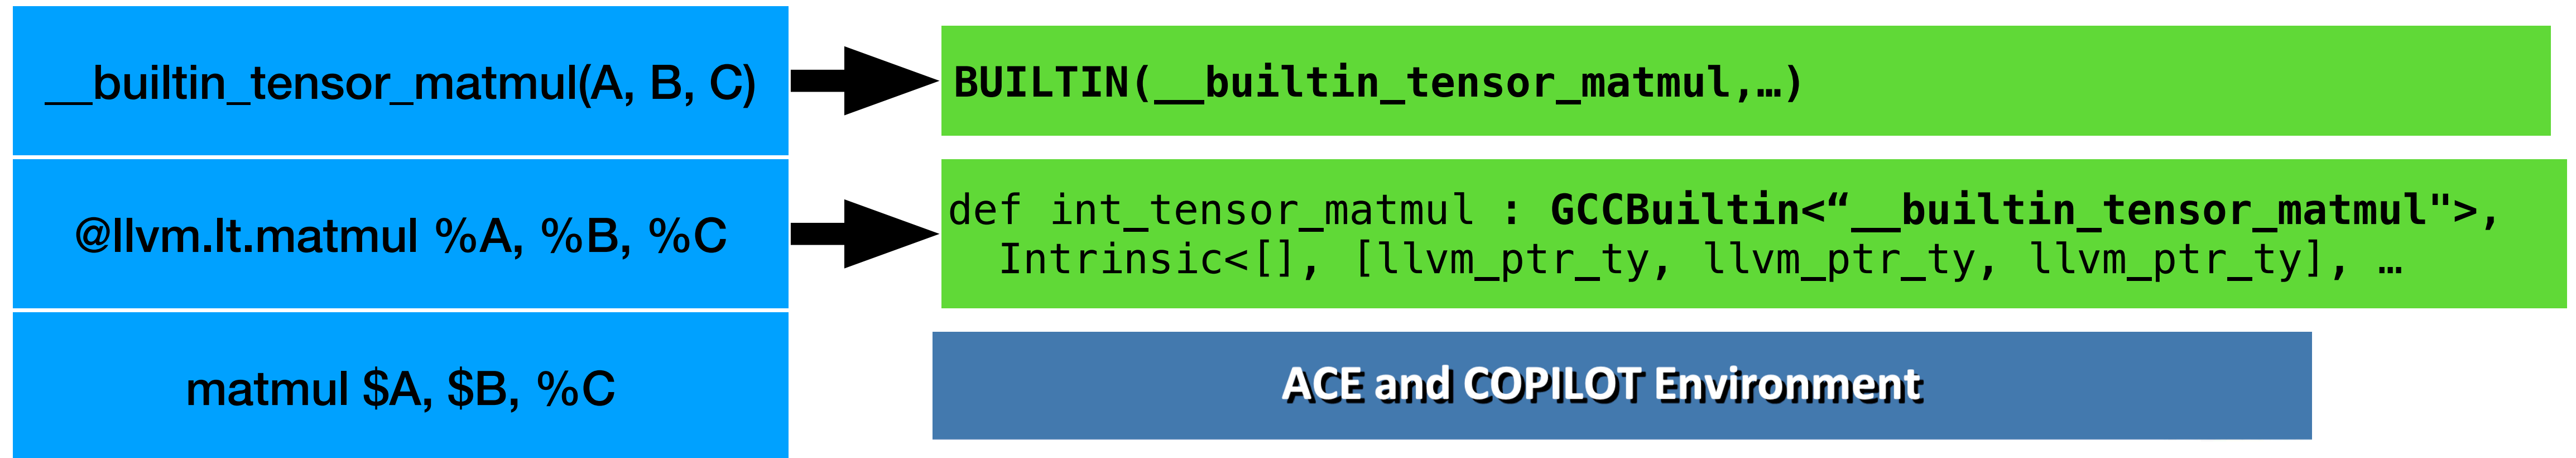
- From 128 \rightarrow 14 cores + 1 oMac
- Core: Andes NV27 + LT's NPU : **coprocessor** (no PC: Program Counter)
- RISCV toolchain
- No Atomic, **mailbox** instead



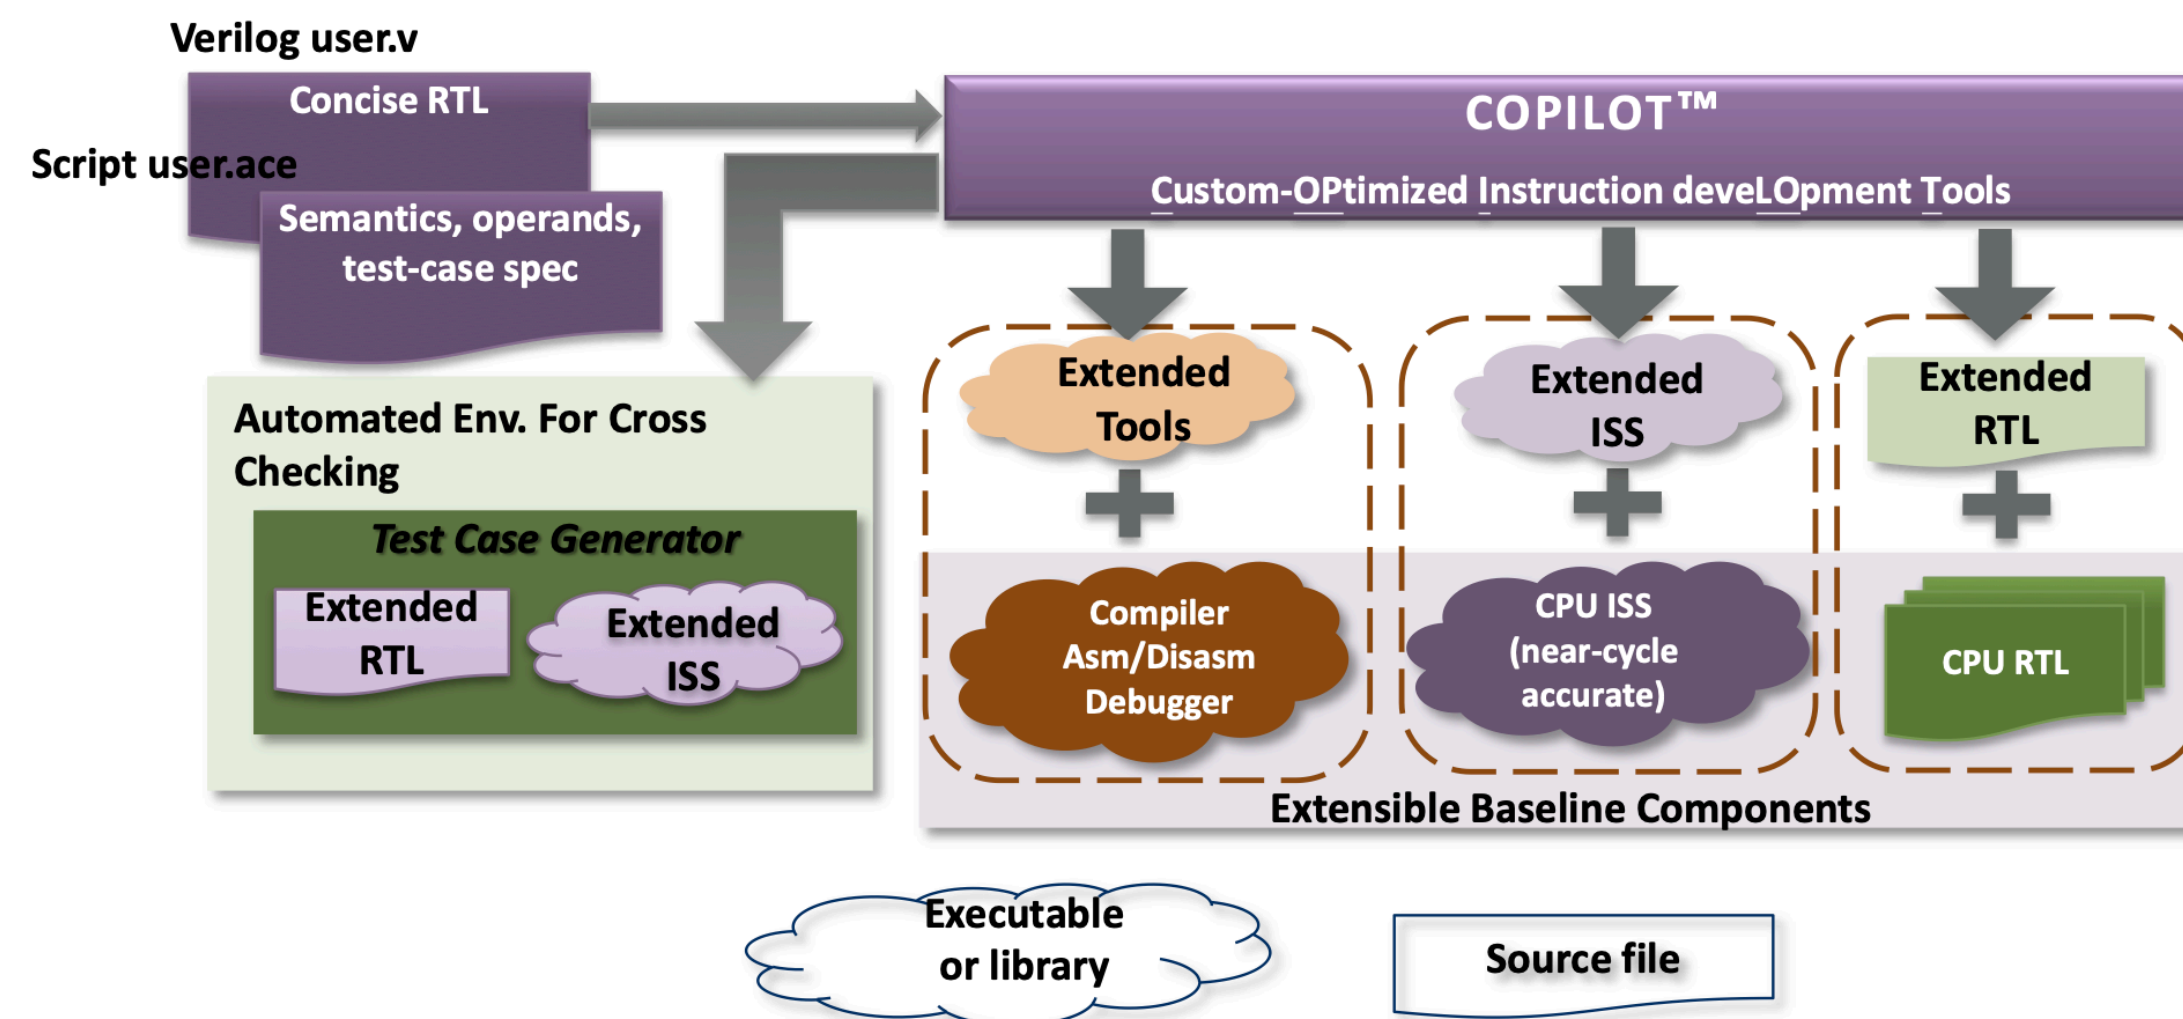
2D Torus illustration

C++ Compiler

`__builtin(clang) + llvm-intrinsic(llvm) + npu.td —> CodeGen`

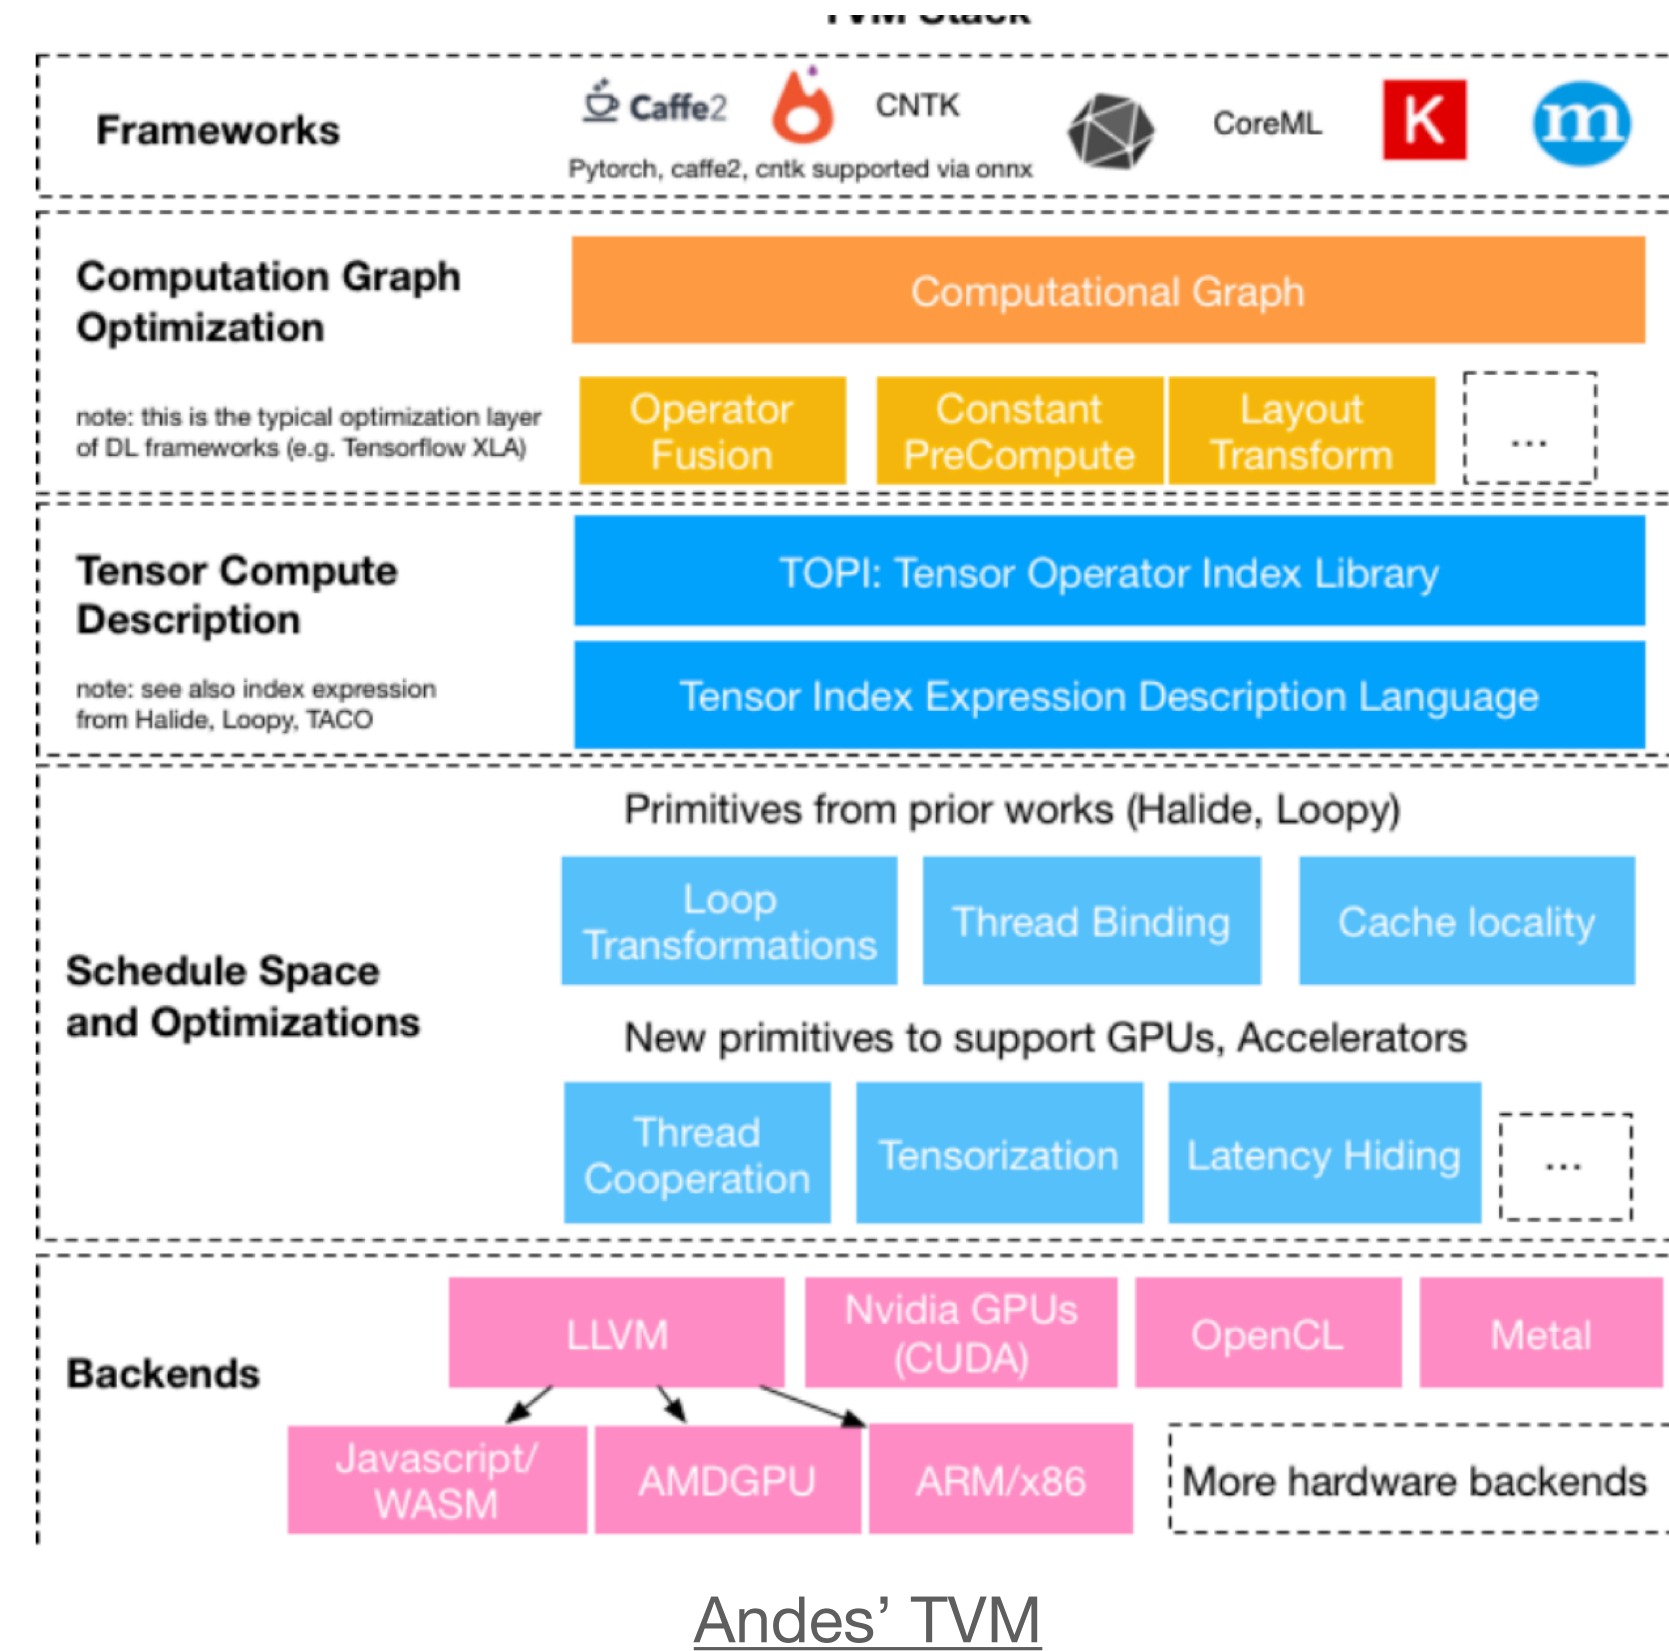

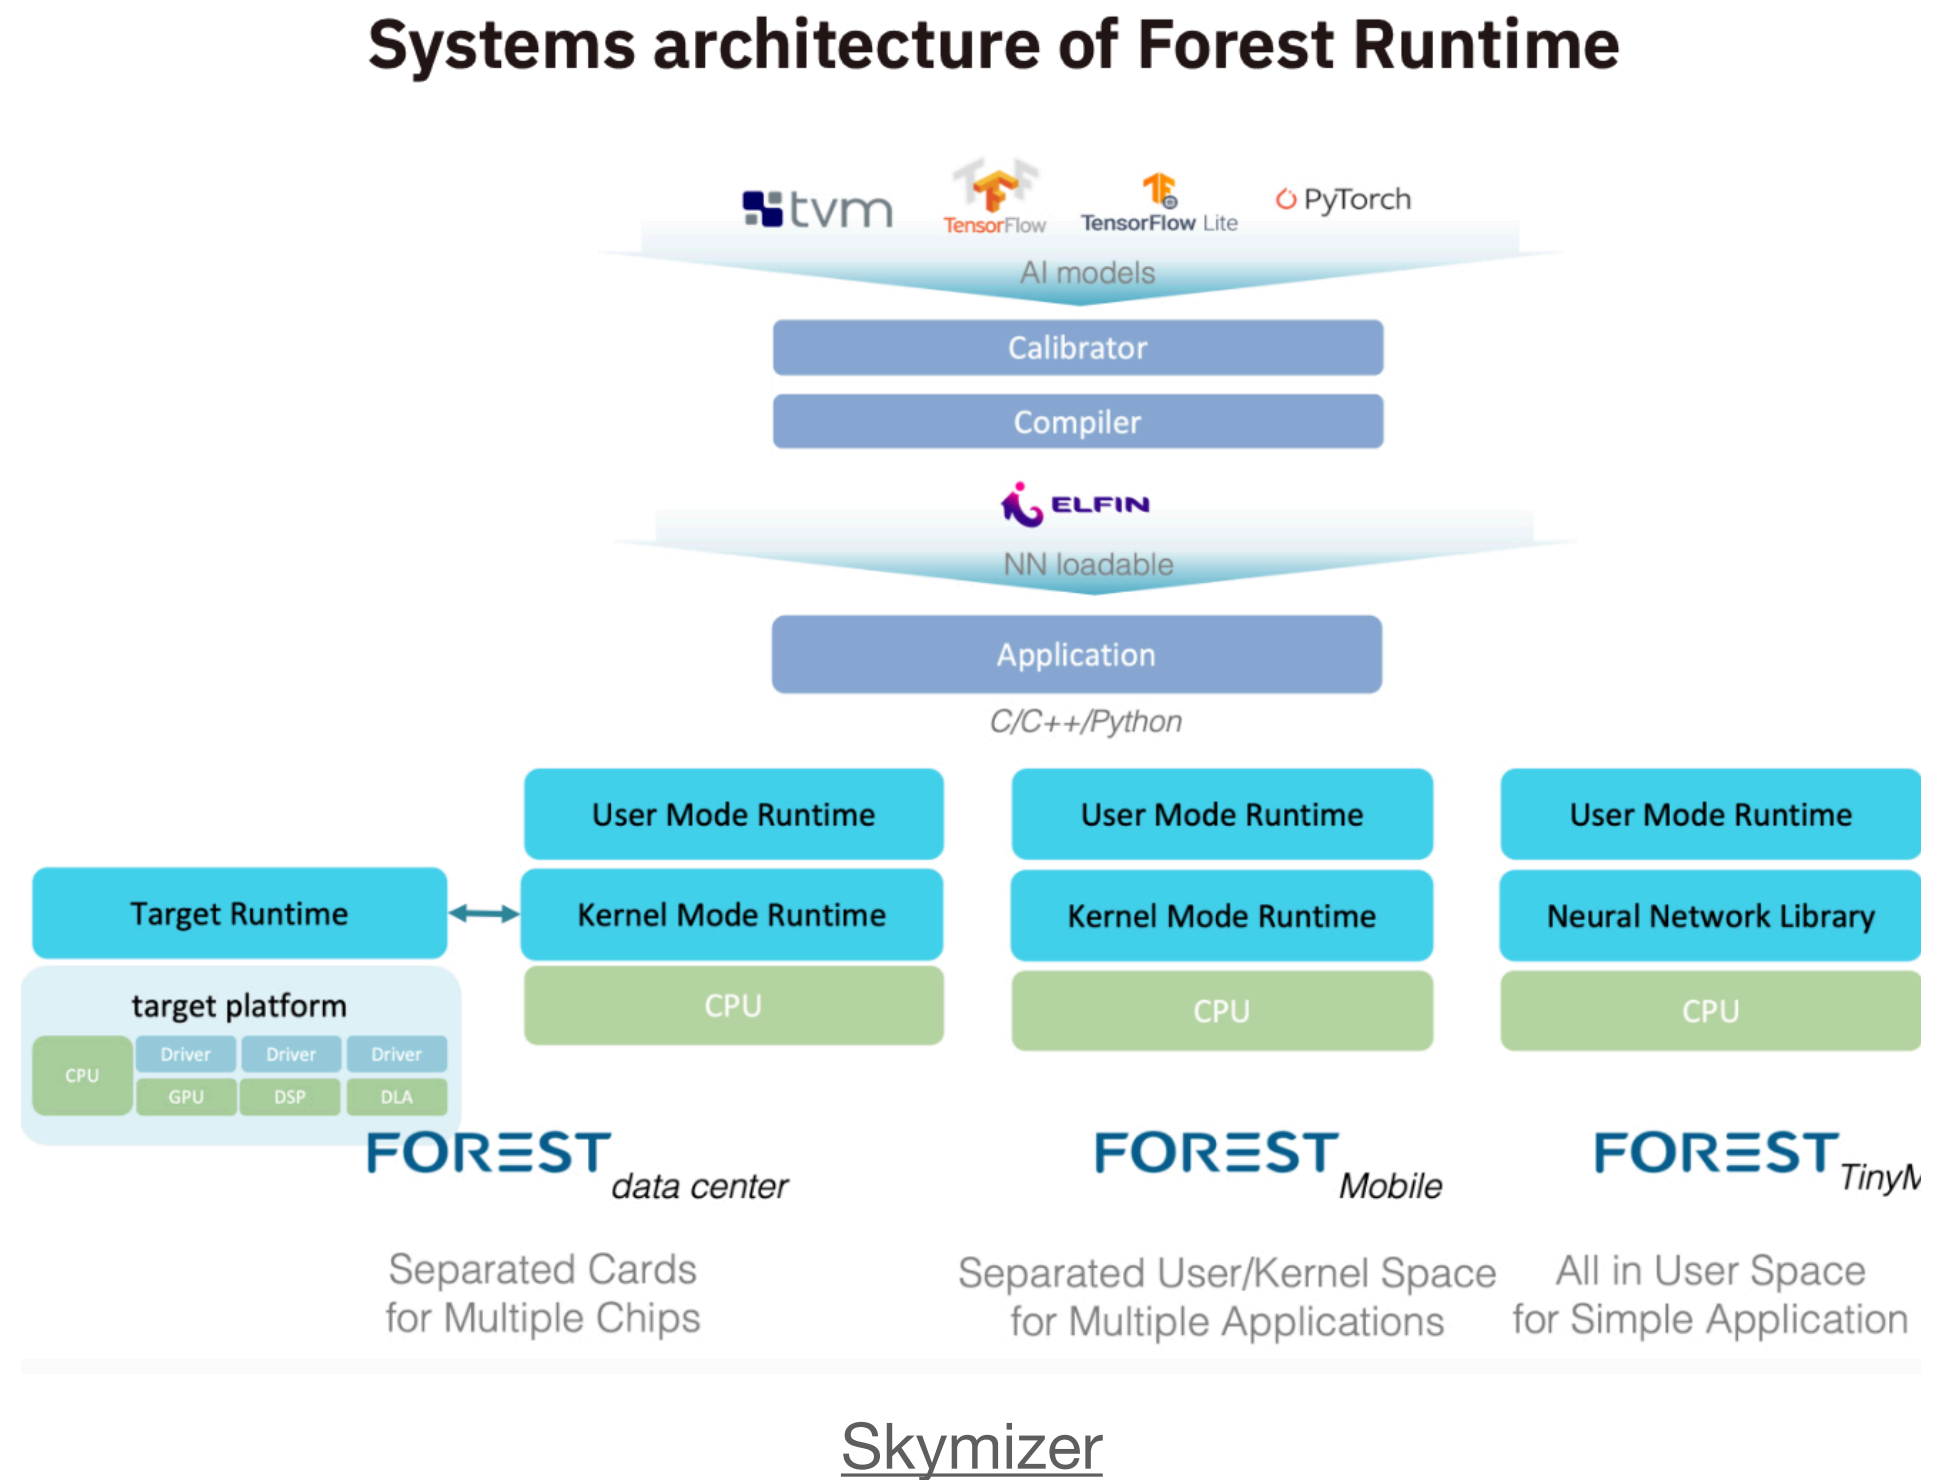


Conv, max_pool, ... , load, store, and IO-control.
30 instructions to implemented



AI SW — Plan for 128-cores

Skymizer + Andes' TVM



TVM for RISC-V Architecture

Andes' TVM for RISCV

But China-Startup's Funding



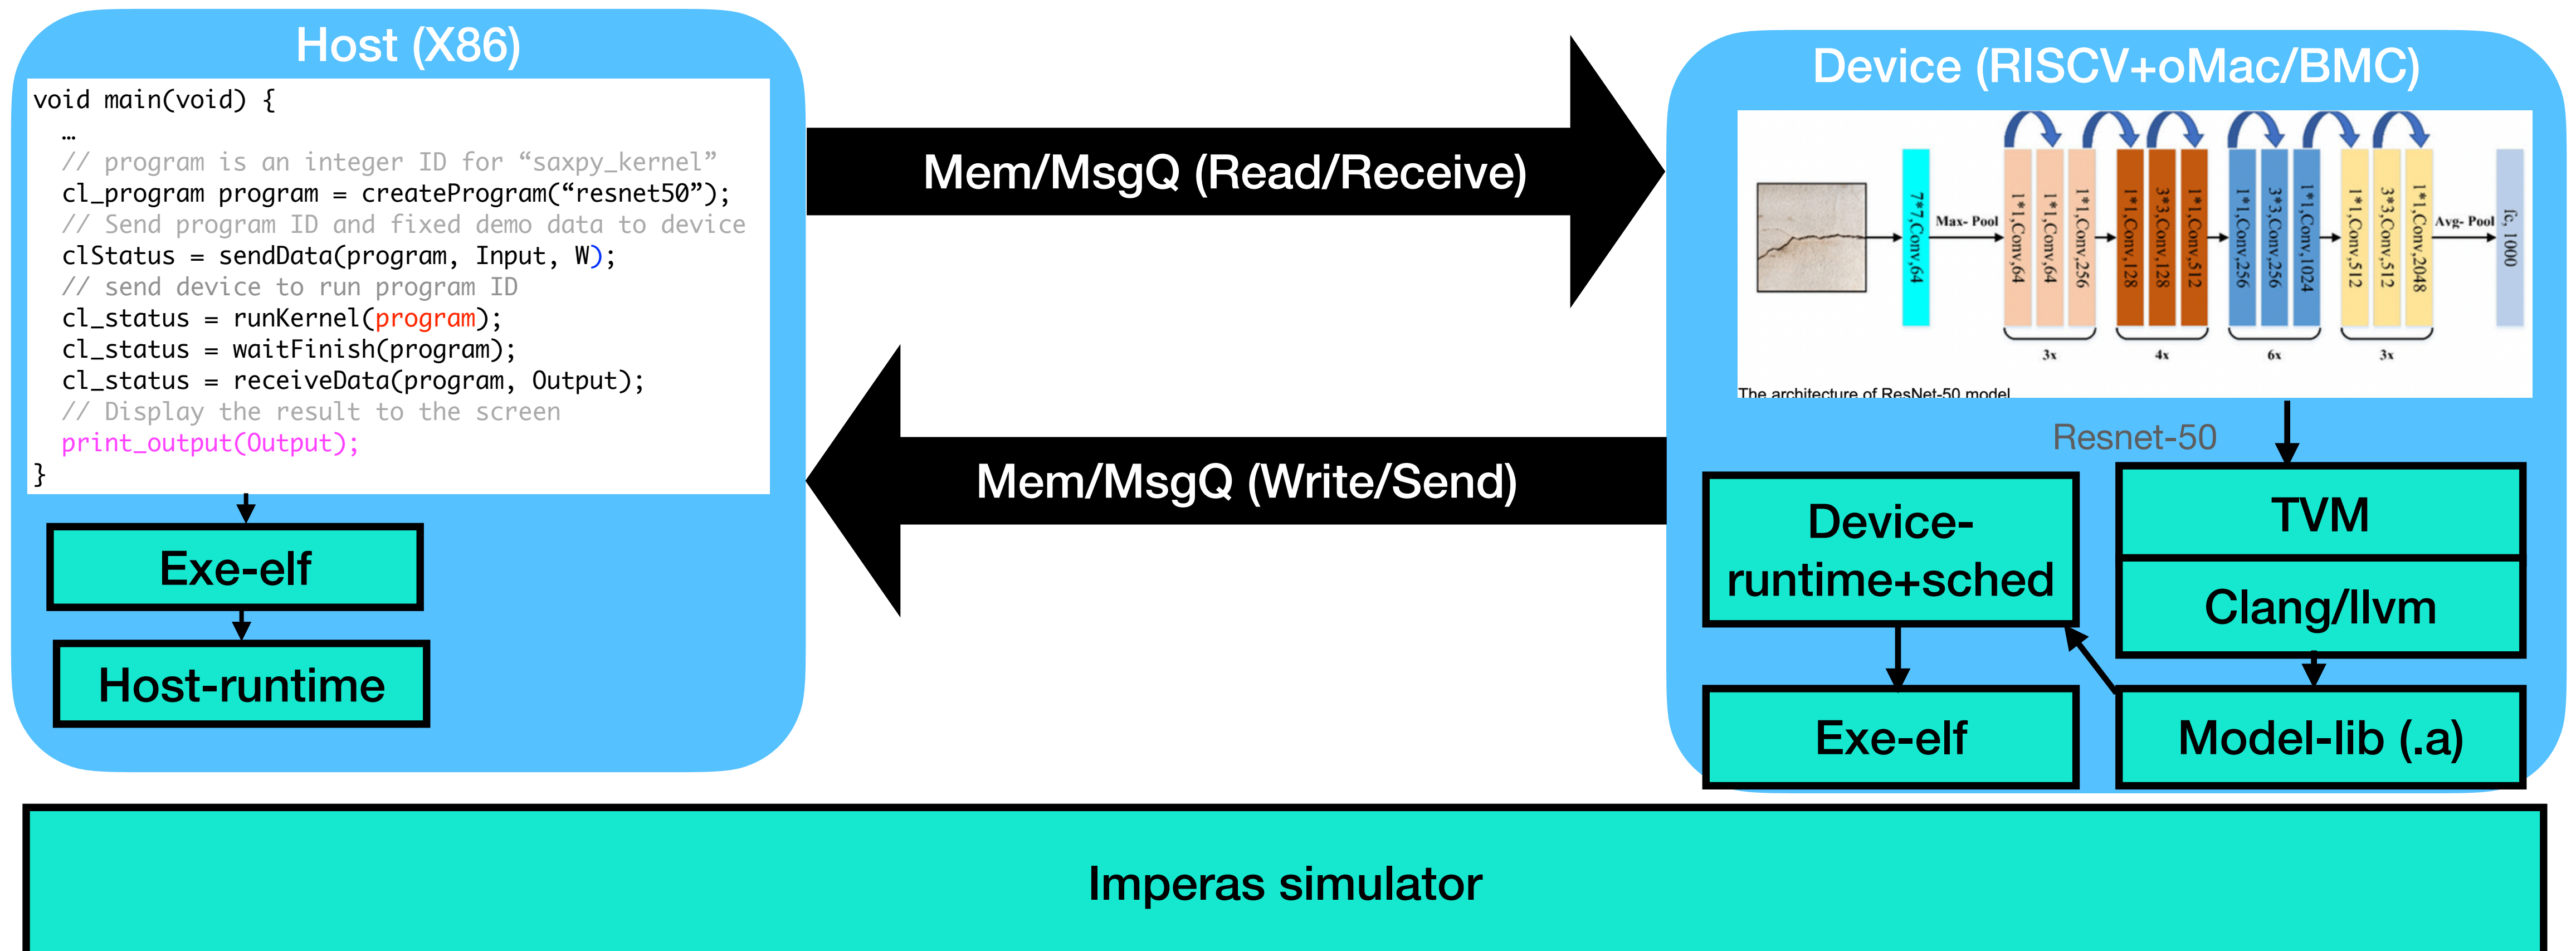
128-Cores

14-Cores

The diagram consists of a thick black arrow pointing downwards and to the right. Above the arrow's starting point is the text '128-Cores' in a bold, dark gray font. At the arrow's tip is the text '14-Cores' in the same font style.

AI (Resnet50, SSD, Monte Carlo) - stage 1

Effort: [Andes' TVM] or ONNX2CApi's compiler



Data science applications - stage 2

Porting OpenCL (pocl) to RISCv

```
void main(void) {
```

```
...
// Create memory buffers on the device for each vector
cl_mem A_clmem = clCreateBuffer(context, CL_MEM_READ_ONLY, VECTOR_SIZE * sizeof(float), NULL, &clStatus);
...
cl_mem B_clmem = clCreateBuffer(context, CL_MEM_READ_ONLY, VECTOR_SIZE * sizeof(float), NULL, &clStatus);
cl_mem C_clmem = clCreateBuffer(context, CL_MEM_WRITE_ONLY, VECTOR_SIZE * sizeof(float), NULL, &clStatus);
```

```
// Copy the Buffer A and B to the device
clStatus = clEnqueueWriteBuffer(command_queue, A_clmem, CL_TRUE, 0, VECTOR_SIZE * sizeof(float), A, 0, NULL, NULL);
clStatus = clEnqueueWriteBuffer(command_queue, B_clmem, CL_TRUE, 0, VECTOR_SIZE * sizeof(float), B, 0, NULL, NULL);
```

```
// Create a program from the kernel source
cl_program program = clCreateProgramWithSource(context, 1, (const char **)&saxpy_kernel, NULL, &clStatus);
```

```
// Build the program
clStatus = clBuildProgram(program, 1, device_list, NULL, NULL, NULL);
```

```
// Create the OpenCL kernel
cl_kernel kernel = clCreateKernel(program, "saxpy_kernel", &clStatus);
```

```
// Set the arguments of the kernel
clStatus = clSetKernelArg(kernel, 0, sizeof(float), (void *)&alpha);
```

```
...
// Display the result to the screen
for(i = 0; i < VECTOR_SIZE; i++)
    printf("%f * %f + %f = %f\n", alpha, A[i], B[i], C[i]);
```

```
// Finally release all OpenCL allocated objects and host buffers.
clStatus = clReleaseKernel(kernel);
```

```
...
}
```

Host (X86)

Device (RISCv+oMac)

Mem/MsgQ

Mem/MsgQ

```
//OpenCL kernel which is run for every work item created.
const char *saxpy_kernel =
"__kernel                                \n"
"void saxpy_kernel(float alpha,          \n"
"    __global float *A,                  \n"
"    __global float *B,                  \n"
"    __global float *C)                  \n"
"{                                         \n"
"    //Get the index of the work-item    \n"
"    int index = get_global_id(0);        \n"
"    C[index] = alpha* A[index] + B[index]; \n"
"}                                         \n";
```


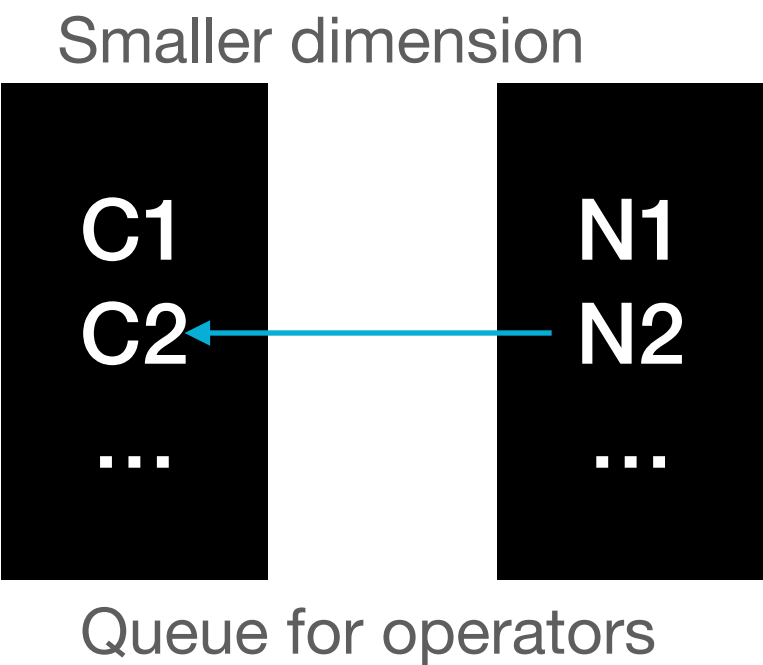
Device Runtime Scheduler for stage 1

Polling

Runtime scheduling -> solve unknown time for DRAM-SRAM

```
for (;;) {
    dispatch(queues); // select ready operator from queue
    run_job(); // operator
    update(queues);
}

// for instance:
Void Addv() {
    for (Ai)
        addv(Aij);
}
```




- Disadvantage:
- Delay BMC-issue after finish RISC-V op.
- Advantage:
- Save ~300 cycles for context switch (registers save & restore)

Device Runtime Scheduler for stage 1

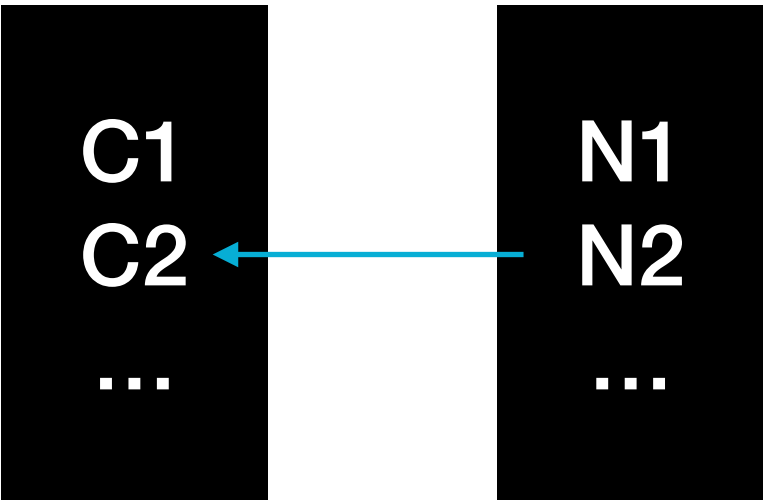
Skymizer: ISR (Interrupt Service Routine)

```
for (;;) {
    dispatch(queues);
    run_job_and_update(queues); // operator
}
```

```
// for instance:
Void Addv() {
    for (Ai)
        addv(Aij); // <- ISR(OMAC finish)
}
```



Context switch:
~300 cycles



Queue for ready operators

- Disadvantage:**
- Cycles of processing queues are extra costs (~300 cycles)
 - Doable when BMC takes ~3,000 (matmul 64*64 : assume 50 cycles)
 - **4MB/4KB=1000, 1000*50=50,000: doable**
- Advantage:**
- Issue BMC as soon as possible
 - Connect to ONNC(Skymizer) for Vanguard

LT's TaskGraph

a.k.a CudaGraph, TaskGraph(DPC++/SYCL/OpenCL)

- MyBookSection
- CodeGen information working with runtime for scheduling.

```
void task_graph() {  
    // define nodes, dependencies and params  
    graph_t A, B, C, D;  
    Type1 X;  
    graph_t d_B[] = {A}, d_C[] = {A}, d_D[] = {B, C};  
  
    createGraphNode(&A, X, 0, 0);  
    createGraphNode(&B, X, d_B, 1);  
    createGraphNode(&C, X, d_C, 1);  
    createGraphNode(&D, X, d_D, 2);  
}
```

